

SEO + AEO — Skill Source Document

Instructions for the Claude Code session reading this file

You are being handed this document to **build a Claude Skill from it**, not to summarize it for a human reader. Treat it as source material for skill authoring, at full depth — do not compress, simplify, or "teach" it back.

Do this:

1. **Create a new skill** (suggested name: `seo-aeo`, or `seo-aeo-audit` if you want to split it into two — your call, but a single skill covering both authoring guidance and the audit checklist is the more useful shape for one agent to load and act on end-to-end).
2. Use the **skill-creator** skill if it's available in this session (check first) to scaffold and validate the skill package properly — correct `SKILL.md` frontmatter, folder layout, and any `references/` subfiles the skill-creator convention expects.
3. **Preserve every source-confidence label** (`[OFFICIAL]` , `[CONSENSUS]` , `[UNCERTAIN]`) from Part 1 — they are load-bearing. A skill that drops them and states everything as fact is a *worse*, not simplified, version of this document. Keep citation URLs.
4. **Keep the audit checklist in Part 2 as a literal, checkable checklist** inside the skill (not prose) — its entire value is that each item is phrased so a page can be graded PASS/FAIL/N/A against it. This is meant to be the skill's primary *action* — "audit this page/URL against seo-aeo" should be something the skill can actually walk through mechanically.
5. The skill should trigger when a user is: writing or reviewing content meant to rank in Google and/or get cited by AI answer engines (ChatGPT, Perplexity, Google AI Overviews); implementing technical SEO (sitemap, robots.txt, canonical, structured data); checking Core Web Vitals; choosing what Schema.org types to add; or auditing a specific page/URL for SEO+AEO readiness.
6. Do not invent claims beyond what's here. If a user asks something this document doesn't cover, the skill should say so rather than guess — that discipline is the whole point of the confidence labels.
7. This document is intentionally topic-agnostic (no specific website, niche, or keyword baked in) — keep the resulting skill that way too, so it's reusable across any project, not just the one it was written for.

Everything below this line (Part 1 and Part 2) is the actual source content, unedited from the research and verification work it came from.

Part 1 — Reference

This is the authoritative, cited reference the demo site must satisfy. It is **topic-agnostic**: every principle applies to any content site, whatever the final subject.

How to read the source labels. Every non-obvious claim is tagged:

- **[OFFICIAL]** — stated by Google, Schema.org, or web.dev in their own documentation. Treat as ground truth.
- **[CONSENSUS]** — not an official rule, but well-established agreement among serious practitioners and repeatedly observed. Reliable, but not a guarantee.
- **[UNCERTAIN]** — genuinely unprovable or contested. Do not build a plan that *depends* on it.

The single most important honesty caveat, up front. Google states plainly that following its guidance does **not** guarantee crawling, indexing, or ranking: *"Google doesn't guarantee that it will crawl, index, or serve your page, even if your page follows the Google Search Essentials."* **[OFFICIAL]** ([Search Essentials](#)).

And there is **no guaranteed formula that forces an AI engine to cite you**. Anyone who promises a specific keyword ranking or a guaranteed AI citation by a date is selling certainty that does not exist. What we *can* do is remove every technical reason to be excluded, and maximize the odds — which on a low-competition long-tail term is genuinely good.

Hebrew note (הערה בעברית): ככה הם מופיעים בכל הכלים (הערה בעברית): המונחים באנגלית בכוונה — מוסיף פה ושם גלוסה קצרה בעברית כדי לקבע את המושג. (Search Console, web.dev).

1. Crawl / Index fundamentals (הבסיס — איך גוגל מוצא ומאחסן דפים)

Google Search works in three stages, and a page must clear all three before it can rank **[OFFICIAL]**

([How Search Works](#)):

1. **Crawling** (סריקה) — Googlebot discovers URLs (from links, sitemaps, and previously known pages) and downloads text, images, and video. Google does not accept payment to crawl more often. **[OFFICIAL]**

2. **Indexing (אינדוקס)** — Google analyzes the downloaded content, processes the page (including running its JavaScript, see §3), and stores it in the index. During indexing Google groups duplicate/near-duplicate pages and picks one **canonical** version to represent the cluster. **[OFFICIAL]**
3. **Serving / Ranking (דירוג)** — for a given query, Google returns what its ranking systems judge the most relevant, highest-quality results for that user, location, device, and language. **[OFFICIAL]**

Hard technical gate: Google only indexes pages served with an **HTTP 200 (success)** status code. Client-error (4xx) and server-error (5xx) pages are not indexed. **[OFFICIAL]**
([Technical requirements](#)).

Practical implications for the build:

- If a page is not crawled → it cannot be indexed → it cannot rank. Crawl is the foundation; fix crawl problems before touching content.
 - `robots.txt` controls **crawling**, not indexing. A URL blocked in `robots.txt` can still appear in Search (as a bare URL with no snippet) if other pages link to it — and because Googlebot never fetched it, it will **never see a `noindex` tag on that page**. To keep a page out of the index, allow crawling and use a `noindex` meta robots tag (or `X-Robots-Tag` header), **not** a `robots.txt` disallow. **[OFFICIAL]**
([Block indexing with noindex; robots.txt intro](#)).
-

2. On-page & content — E-E-A-T, search intent, answer-first structure

2.1 People-first, helpful content

Google's ranking systems aim to reward **original, high-quality, people-first** content and to demote content made primarily to manipulate rankings. The "helpful content" approach began in 2022 as a signal and was folded into Google's **core ranking systems in the March 2024 core update**. **[OFFICIAL]**
([Creating helpful content](#)).

Google publishes **self-assessment questions** to judge your own content. The useful ones to hold the build to **[OFFICIAL]** (same page):

- Does the content provide original information, reporting, research, or analysis — not just rehash what other sites already say?
- Is it written by, or does it clearly demonstrate, **first-hand expertise or experience**?

- After reading, will someone feel they learned enough to accomplish their goal, or do they need to search again?
- Would you trust this content for a decision that matters (money, health, safety)?
- Is it free of obvious errors, and does it read like it was made for people rather than to hit a keyword?

2.2 E-E-A-T (Experience, Expertise, Authoritativeness, Trust)

E-E-A-T is **not a single direct ranking score** — it is the *framework* Google's quality raters and systems use to judge whether content is reliable.

[OFFICIAL] Google states: "*Trust is the most important member of the E-E-A-T family... the other members of E-E-A-T contribute to trust, but content isn't always trustworthy if it shows those aspects.*"

([Helpful content page, E-E-A-T section](#); [Search Quality Rater Guidelines PDF](#)).

The "**who / how / why**" test Google recommends **[OFFICIAL]**:

- **Who** created the content — is there a real, credited author with a bio and credentials? (Especially important here: the owner's blindness gives genuine *first-hand Experience* on accessibility topics — this is a real, hard-to-fake E-E-A-T advantage. Make the author identity explicit.)
- **How** it was created — disclose method, and any AI involvement where it would matter to the reader.
- **Why** it exists — to help people, not primarily to rank or to serve ads.

On AI-generated content: Google does **not** ban AI content. Its stance is that it rewards high-quality content *however it is produced*, but that using automation to generate content *primarily to manipulate rankings* violates spam policies. **[OFFICIAL]**

([Google Search's guidance about AI-generated content](#)).

Practical rule for this project: AI can draft and assist, but a human with real expertise must direct, verify, and own the result.

2.3 Search intent (כוונת החיפוש)

Ranking is intent-matching, not keyword-matching. A page ranks when it satisfies the *task behind the query* better than alternatives. **[CONSENSUS]** The four classic intent types: **informational, navigational, commercial, transactional**. Before writing, look at what already ranks for the target query — the current page-1 results reveal the intent Google has decided the query has. Match that format (guide vs. list vs. definition vs. tool). **[CONSENSUS]**

2.4 Answer-first structure ("מבנה "תשובה קודם")

Structure that helps both human skimmers and AI extractors **[CONSENSUS]**:

- Put a **direct, self-contained answer in the first 1–3 sentences** under each heading — before the background and caveats. AI answer engines and Google's featured snippets pull short, standalone answer passages; buried answers get passed over.
- Use a **clear heading hierarchy** (one `<h1>` , descriptive `<h2>` / `<h3>`) that mirrors the questions a reader would ask. Phrase headings as the actual questions where natural.
- Prefer **short paragraphs, lists, and tables** for facts, steps, and comparisons — these are the chunks retrieval systems extract cleanly.
- One primary topic per page (**topical focus**). Don't dilute a page across unrelated subjects.
- Include the **question and its key entities in the visible text**, not only in a title tag.

3. Technical SEO

Area	What to do	Source / label
HTTPS	Serve the whole site over HTTPS with a valid certificate; HTTPS is a lightweight ranking signal and a baseline expectation.	OFFICIAL — HTTPS
Status codes	Only 200 pages get indexed. Use 301 for permanent redirects, 404/410 for gone pages, avoid 5xx and "soft 404s" (a 200 page that says "not found").	OFFICIAL
Canonical	When the same/near-same content lives at multiple URLs, pick one canonical and mark it with <code>rel="canonical"</code> . Google treats it as a strong hint, not a command; keep signals consistent (canonical, internal links, sitemap, redirects all pointing the same way).	OFFICIAL — canonicalization
XML sitemap	Publish a sitemap listing your canonical, indexable URLs; reference it in <code>robots.txt</code> and submit it in Search Console. Helps discovery, especially for new sites with few backlinks.	OFFICIAL — sitemaps
robots.txt	Use to manage crawl, never as an indexing control (see §1). Do not block CSS/JS/image resources Googlebot needs to render the page.	OFFICIAL
Mobile	Google indexes with mobile-first crawling — Googlebot Smartphone is the source of truth. Content, links, and meta robots tags must be identical on mobile and desktop; a <code>noindex</code> / <code>nofollow</code> that appears only on mobile can drop the page.	OFFICIAL — mobile-first
JS rendering	Google renders pages by executing JS in an evergreen Chromium Web Rendering Service. It works, but rendering is a second, deferred pass and can fail or lag. Ensure primary content and links exist in the rendered DOM; prefer server-side rendering / static generation / hydration so content does not depend on a click or a slow client fetch. Test with the URL Inspection tool's rendered HTML.	OFFICIAL — JS SEO basics

Meta essentials (title & description): the `<title>` and meta description do not directly rank the page, but they are what Google usually shows in the SERP snippet and heavily influence **click-through rate**. Write a unique, accurate, descriptive title per page; keep the primary query near the front. **[CONSENSUS]** ([Control your title links; snippets/meta description](#)).

4. Core Web Vitals (CWV) — current metrics and thresholds

Core Web Vitals are three real-user (field) metrics for loading, interactivity, and visual stability. A page/URL group **passes** only when **all three** meet the "good" threshold at the **75th percentile** of real visits (i.e. 75% of visits are at least this good), assessed separately for mobile and desktop.

[OFFICIAL]

([web.dev — Core Web Vitals; thresholds definition](#)).

Metric	Measures	Good	Needs improvement	Poor
LCP — Largest Contentful Paint	Loading: time until the largest content element renders	≤ 2.5 s	≤ 4.0 s	> 4.0 s
INP — Interaction to Next Paint	Responsiveness: worst-ish latency from a user interaction to the next paint, across the visit	≤ 200 ms	≤ 500 ms	> 500 ms
CLS — Cumulative Layout Shift	Visual stability: how much visible content shifts unexpectedly	≤ 0.1	≤ 0.25	> 0.25

Thresholds per [web.dev](#). **[OFFICIAL]**

INP replaced FID. As of **March 12, 2024**, INP became the official responsiveness Core Web Vital, replacing First Input Delay (FID), which is now deprecated. INP is stricter: it measures *all* interactions in a visit, not just the first, and includes processing + rendering time, not just input delay.

[OFFICIAL]

([web.dev — INP is now a Core Web Vital; Google Search blog](#)).

How Google uses CWV in ranking (be precise): page experience, including CWV, is a ranking signal, but Google is explicit that it is **not a "boost" and does not override relevance** — great content on a slightly-slower page can still outrank thin content on a fast page. CWV is best treated as a **tiebreaker and a must-not-fail hygiene bar**, and a real UX benefit in its own right. **[OFFICIAL]** ([Understanding page experience](#)).

Field vs. lab: CWV assessment uses **field data** (real users, reported via the Chrome UX Report / CrUX and shown in Search Console's Core Web Vitals report). Lab tools (Lighthouse, PageSpeed Insights lab mode) are for debugging and give a *lab* Performance score — a good Lighthouse score is not the same as passing CWV in the field. **[OFFICIAL]** (web.dev).

Build targets to bake in from day one: server-rendered/static HTML, a compressed and correctly-sized LCP image with `width / height` set, `font-display: swap`, reserved space for any dynamic/embedded content (prevents CLS), and minimal blocking JavaScript (protects INP).

5. Structured data (Schema.org / JSON-LD)

What it is and is not. Structured data is machine-readable markup describing the page's meaning. Google's **recommended format is JSON-LD** (a `<script type="application/ld+json">` block), using the **Schema.org** vocabulary.

[OFFICIAL]

([Intro to structured data](#); [Schema.org](https://schema.org/)).

It does not directly boost rankings. Its job is **eligibility for rich results** (enhanced SERP appearances) and helping machines understand entities and relationships on the page. Correct markup makes you *eligible*; Google still decides whether to show a rich result based on quality and its own rules.

[OFFICIAL]

([Structured data general guidelines](#)).

Two firm rules: (1) mark up only content that is **actually visible on the page** — invisible or fabricated markup is a spam violation; (2) supply all **required** properties for a type or it won't be eligible. **[OFFICIAL]**

Types worth implementing for a content/article site (validate every one in the [Rich Results Test](#) and monitor in Search Console):

- **Organization** (or **Person** for a solo author-site) — site-wide identity: name, logo, `sameAs` links to real profiles. Feeds Google's entity understanding and Knowledge Panel eligibility. **[OFFICIAL]** ([Organization](#)).
- **WebSite** with `SearchAction` (sitelinks search box) — optional. **[OFFICIAL]**
- **Article** / **BlogPosting** / **NewsArticle** — for each content page: `headline`, `author` (link to a **Person** with credentials — reinforces E-E-A-T), `datePublished`, `dateModified`, `image`, `publisher`. **[OFFICIAL]** ([Article](#)).

- **BreadcrumbList** — powers breadcrumb display in results and clarifies site hierarchy. **[OFFICIAL]** ([Breadcrumb](#)).
- **Person** (author entity) — bio, **jobTitle**, **sameAs**. Ties content to a credible, real author. **[OFFICIAL / CONSENSUS]**

Types to know are restricted or gone — do NOT rely on these for rich results:

- **FAQPage** — FAQ rich results were **restricted in August 2023** to well-known government/health sites, and **fully deprecated (removed for all sites) in 2026**. The markup is still a valid Schema.org type and can still aid machine understanding, but it will **not** earn a SERP rich result for a normal site — do not add it expecting extra SERP real estate. **[OFFICIAL]** ([HowTo/FAQ changes](#)).
- **HowTo** — HowTo rich results were **deprecated in 2023** (mobile then desktop). Same story: valid vocabulary, no rich result. **[OFFICIAL]** (same source).

Bottom line: implement **Organization / Person**, **Article**, and **BreadcrumbList** cleanly and keep them valid; treat FAQ/HowTo markup as optional semantic context only, never as a rich-result strategy.

6. Measurement — Google Search Console (GSC) signals to watch

GSC is the **primary, authoritative** feedback loop (it reports Google's own data, unlike third-party estimators). The proof of "first-page ranking" the project needs comes from here. **[OFFICIAL]** ([Search Console help](#)).

Indexing / Pages report — is the page even in the index?

- Confirm target URLs show **"Indexed"** (Page indexing report). Investigate any "Discovered – not indexed" / "Crawled – not indexed" / "Excluded" states first — an un-indexed page can never rank. **[OFFICIAL]**
- Use **URL Inspection** to see live crawl status, the **rendered** HTML (verifies JS content is seen), canonical Google chose, and to **Request Indexing** for a new/updated page. **[OFFICIAL]**

Performance report — four metrics, all from Google **[OFFICIAL]**:

- **Impressions** — how often your URL appeared in results.
- **Clicks** — how often users clicked it.
- **CTR** — clicks ÷ impressions (a snippet-quality signal; improve via title/description).
- **Average position** — mean ranking position across impressions. **"First page"**

≈ **average position** ≤ ~10 for the **target query** — this is the number to screenshot as proof. Filter Performance by the exact query to show it.

Also watch: the **Core Web Vitals** report (field pass/fail by URL group, mobile + desktop) and **Mobile Usability**/page-experience signals; the **Sitemaps** report (submitted vs. indexed); and any **Manual Actions** / **Security** issues (must be zero). **[OFFICIAL]**

Reading the funnel: low impressions on a page you expected to rank → an *upstream* problem (indexing, crawlability, or relevance/intent mismatch), not a CTR problem. Fix in that order. **[CONSENSUS]**

7. AEO — Answer Engine Optimization (getting retrieved and cited by AI)

Honesty first. There is **no guaranteed way to make an AI cite you**. AI answer engines are non-deterministic, their retrieval and ranking are proprietary and change often, and citation of any specific source varies run-to-run. Everything below raises probability; nothing sets it to 1. **[UNCERTAIN as to any specific citation]**

7.1 How the major answer engines actually retrieve and cite

Most AI answer engines use **RAG (Retrieval-Augmented Generation)**: rather than answering purely from the model's memory, the system **retrieves live web documents, then generates an answer grounded in that retrieved text, attaching citations**. **[CONSENSUS — vendor-described]**

- **Google AI Overviews / AI Mode** draw from **the same Google index and core ranking systems** as normal Search. They use a "**query fan-out**" technique — issuing several related sub-queries and synthesizing the results. Google states there are **no special requirements and no special structured data** to appear; a page must simply be **indexed and eligible to show in Search with a snippet**. This means classic SEO is the AEO groundwork for Google. **[OFFICIAL]** ([AI features & your website](#); [Succeeding in AI search](#)).
- **Perplexity** runs real-time web retrieval on **every** query (hybrid keyword + embedding retrieval, then re-ranking), and cites sources inline by default — it is the most citation-dense mainstream engine. **[CONSENSUS — reverse-engineered / vendor statements, not a published spec]**
- **ChatGPT Search** uses web results (historically Bing-indexed) but only invokes web search on a *fraction* of queries; otherwise it answers from

training data. Being in the underlying web index is a prerequisite for being cited when it does search. [CONSENSUS]

The load-bearing takeaway: for every one of these, **being crawlable, indexed, and genuinely useful is the prerequisite**. You cannot be cited from a page the engine cannot retrieve or does not trust.

7.2 What actually raises the odds of being cited (practitioner consensus)

[CONSENSUS] unless marked otherwise:

1. **Be indexable and technically healthy** — everything in §1–§5. For Google's AI features this is [OFFICIAL] the whole requirement.
2. **Answer-first, extractable structure** (§2.4) — a clean, self-contained answer under a clear question-shaped heading is the unit these systems lift.
3. **One clear claim per passage, with specifics** — concrete facts, numbers, dates, named entities. Vague prose is hard to ground and rarely quoted.
4. **Demonstrable E-E-A-T and trust** — real named author, credentials, citations to primary sources, "last updated" dates. Engines and their rankers favor sources they can treat as reliable.
5. **Corroboration across the web** — being mentioned/consistent across multiple independent, reputable sources correlates with being surfaced. (A brand-new site with zero external mentions is at a disadvantage here — expect the AI citation to be the *slower* of the two proof goals.) [CONSENSUS / UNCERTAIN]
6. **Match real question phrasing** — write the way people ask, including natural long-tail questions, so the passage matches the fanned-out sub-queries.
7. **Keep content fresh and accurate** — outdated or wrong facts get filtered out and erode trust.

7.3 Things sold as AEO magic that are NOT proven

- **llms.txt** — a proposed file listing "AI-friendly" content. **No major AI provider (Google, OpenAI, Anthropic, Meta) has committed to reading it in production**; Google has publicly declined to support it. Independent analyses find **no correlation between having llms.txt and being cited more**. It is low-cost and harmless to add, but do **not** count it as a citation lever. [CONSENSUS — strongly evidenced skepticism] ([SE Ranking analysis](#)).
- **Adding FAQPage / HowTo schema for AI** — see §5; these no longer produce rich results and are not a demonstrated citation trigger. Structure the actual Q&A in visible content instead. [CONSENSUS]
- **"Guaranteed AI citation" services / a secret schema that forces citation** — do not exist. [UNCERTAIN → treat as false]

7.4 SEO ↔ AEO relationship

For Google specifically, **AEO is a superset of good SEO, not a separate discipline** — Google says its normal best practices are what make you eligible for AI features **[OFFICIAL]**. For non-Google engines, the *same* fundamentals (retrievable, structured, trustworthy, corroborated) do most of the work **[CONSENSUS]**. So the strategy is one strategy: **build an excellent, technically flawless, genuinely authoritative page, and both goals are served by the same work.**

8. What this means for the demo (the bar to clear)

Mapping back to the project's three proof goals:

1. **First-page Google ranking on a long-tail keyword, shown in Search Console.**
Achievable because (a) long-tail = low competition, and (b) the owner has a real E-E-A-T edge on accessibility topics. Requires: indexable page (§1), intent-matched answer-first content (§2), clean technical SEO (§3), passing CWV (§4). Proof = GSC Performance report, filtered to the query, average position $\leq \sim 10$.
2. **At least one AI-engine citation.** No guarantee (§7). Maximize odds via §7.2. Expect this to lag the Google ranking, because it also benefits from external corroboration a new site has to earn.
3. **Full technical health.** Green CWV in the field (§4), valid structured data in the Rich Results Test (§5), mobile-first-clean and HTTPS (§3), zero manual actions (§6).

Decisions the owner should make (flagged for a human, not for me to decide):

- **The one long-tail keyword** is still unchosen — this gates everything.
Pick a specific low-competition question where the owner's first-hand experience is a real advantage.
 - **Confirm the exact bar with the friend** — "first page" (≤ 10) vs. "top 3", and which AI engine's citation counts.
 - **Author identity / disclosure** — how prominently to surface the owner's real name and lived experience (strong E-E-A-T, but a personal privacy choice for a 15-year-old).
-

Part 2 — Audit Checklist

Use this to audit **one page** objectively. Each item is written so the answer is **PASS / FAIL / N/A**, with *how to check* named. This is the "ranking checker" the owner asked for. Companion reference (the *why* + sources):

[seo-aeo-reference.md](#) .

Scoring rule. Sections A (Crawlability) and D (Performance/CWV) are **gates** — any FAIL there means the page cannot reliably rank, fix first. Sections B, C, E, F are quality multipliers. Aim for **100% of gates + ≥90% of the rest** before calling a page "done."

Tools you'll use (all free):

- Google Search Console → URL Inspection, Performance, Page indexing, Core Web Vitals
 - [Rich Results Test](#)
 - [PageSpeed Insights](#) (field CrUX data + lab)
 - Chrome DevTools (Lighthouse, device toolbar, Elements = rendered DOM)
 - `view-source:` + the live URL; `curl -I <url>` for status/headers
-

A. Crawlability & Indexing (GATE — כניסה לאינדקס)

- [] **A1. Returns HTTP 200.** `curl -I <url>` shows `200` . No 3xx chain to the final URL, no 4xx/5xx. (*Only 200 pages get indexed.*)
- [] **A2. Not blocked in `robots.txt`.** The URL and its CSS/JS/image resources are crawlable. Check `/robots.txt` and URL Inspection → "Crawl allowed? Yes".
- [] **A3. Indexable directives.** Page has **no** `noindex` (meta robots or `X-Robots-Tag` header). Verify in rendered HTML *and* response headers, not just source.
- [] **A4. Self-referencing / correct canonical.** `<link rel="canonical">` points to this page's preferred URL (or the intended canonical). Google's chosen canonical (URL Inspection) matches your intent.
- [] **A5. In an XML sitemap** that is submitted in Search Console, and the sitemap lists the canonical URL (not a redirecting/duplicate variant).
- [] **A6. Actually indexed.** URL Inspection / Page indexing report says **"URL is on Google" / Indexed** — not "Discovered/Crawled – currently not indexed" or "Excluded".
- [] **A7. Reachable by internal links.** At least one crawlable `<a href>` from another indexed page points here (not orphaned, not JS-only navigation).

- [] **A8. Content present in rendered DOM.** URL Inspection → "View rendered HTML" (or DevTools Elements) shows the main text and links — i.e. content does not depend on a user action or a failed client fetch.

B. Content & Search Intent (תוכן וכוונת חיפוש)

- [] **B1. One primary query/topic** per page; the page's purpose is obvious.
- [] **B2. Intent match.** The page format (guide / list / definition / tool) matches what already ranks on page 1 for the target query. (*Check by searching the query and eyeballing the top results.*)
- [] **B3. Answer-first.** A direct, self-contained answer appears in the first 1–3 sentences (and under each major heading) before background/caveats.
- [] **B4. Heading hierarchy.** Exactly one `<h1>`; logical `<h2>` / `<h3>`; headings phrased as the questions a reader would ask where natural.
- [] **B5. Original value.** Contains first-hand experience, original analysis, specifics (numbers, dates, named entities) — not a rehash of other pages.
- [] **B6. Named, credible author.** Visible author with bio/credentials; a "last updated" date; sources cited for factual claims. (*E-E-A-T / trust.*)
- [] **B7. Complete.** A reader finishes without needing to search again for the same task.
- [] **B8. Unique, descriptive `<title>`** (primary query near the front) and a compelling meta description. Not duplicated across pages.
- [] **B9. Descriptive URL slug** (words, hyphens, no junk parameters).
- [] **B10. Accurate & current.** No factual errors; no stale claims. For YMYL-ish topics (health/safety/money) accuracy is non-negotiable.

C. Structured Data (נתונים מובנים)

- [] **C1. JSON-LD present** (`<script type="application/ld+json">`), using Schema.org vocabulary. (*JSON-LD is Google's recommended format.*)
- [] **C2. Validates** in the Rich Results Test with **zero errors** (warnings acceptable). No "invalid"/missing required properties.
- [] **C3. Article / BlogPosting** on content pages with `headline`, `author` (→ `Person`), `datePublished`, `dateModified`, `image`, `publisher`.
- [] **C4. Organization or Person** site identity present (name, logo, `sameAs` to real profiles).
- [] **C5. BreadcrumbList** present and matches the visible breadcrumb / site hierarchy.
- [] **C6. Markup mirrors visible content only** — nothing marked up that isn't on the page; no fabricated data. (*Spam-policy violation otherwise.*)

- [] **C7. No reliance on dead rich-result types.** Not counting on `FAQPage` / `HowTo` for a SERP rich result (deprecated/restricted). If used, it's for semantic context only and reflects real on-page Q&A.

D. Performance / Core Web Vitals (GATE — ביצועים)

Judge on field data (PageSpeed Insights "real users" / GSC Core Web Vitals), mobile first. Lab scores are for debugging only.

- [] **D1. LCP ≤ 2.5 s** at the 75th percentile (field). (*Good.*)
- [] **D2. INP ≤ 200 ms** at the 75th percentile (field). (*Good; INP replaced FID in March 2024.*)
- [] **D3. CLS ≤ 0.1** at the 75th percentile (field). (*Good.*)
- [] **D4. Passes on BOTH mobile and desktop** (assessed separately).
- [] **D5. LCP image optimized** — compressed, correctly sized, explicit `width / height`, not lazy-loaded, ideally `preload` ed.
- [] **D6. No layout shift sources** — images/embeds/ads have reserved dimensions; web fonts use `font-display: swap`.
- [] **D7. Minimal blocking JS** — main thread not overloaded (protects INP); defer/async non-critical scripts.
- [] **D8. Lighthouse mobile Performance** used only to *diagnose* D1–D3 (note: a green Lighthouse score $\neq$ passing field CWV).

E. Accessibility & Mobile (נגישות ומובייל — also an E-E-A-T topic-fit edge here)

- [] **E1. Mobile parity** — same content, links, and meta-robots on mobile and desktop (*mobile-first indexing uses the mobile version as source of truth*).
- [] **E2. Responsive, no horizontal scroll**, tap targets adequately sized, legible base font.
- [] **E3. HTTPS** with a valid certificate across the whole page (no mixed content).
- [] **E4. Semantic HTML & landmarks** — real `<h1>` /headings, `<nav>`, `<main>`, lists; proper reading order.
- [] **E5. Images have meaningful alt** text (decorative images `alt=""`).
- [] **E6. Keyboard operable** — all interactive elements reachable/usable by keyboard, visible focus states.
- [] **E7. Sufficient color contrast** (WCAG AA: 4.5:1 body text).
- [] **E8. Accessible names / ARIA** where needed; forms have labels.

- [] **E9. lang attribute set** (`<html lang="...">`), correct for the content language.
- [] **E10. Screen-reader spot check** passes (the owner can do this directly — a genuine authenticity signal for an accessibility site).

F. AEO-readiness (סיכוי להיות מצוטט ע"י מנועי AI — probabilistic, NOT guaranteed)

No item here guarantees a citation. They raise the odds. A hard requirement for Google's AI features is simply: the page passes A–E and is eligible to show with a snippet.

- [] **F1. Retrievable** — passes all of Section A (an AI cannot cite what it can't retrieve/index). *(For Google AI Overviews/AI Mode this is the whole official requirement.)*
 - [] **F2. Extractable answer chunks** — self-contained answers under question-shaped headings (ties to B3/B4); short paragraphs, lists, tables for facts/steps.
 - [] **F3. One clear claim per passage**, with concrete specifics (facts, numbers, dates, named entities) that a model can ground a citation on.
 - [] **F4. Question-phrased content** matching how people actually ask (natural long-tail phrasing appears in visible text).
 - [] **F5. Trust signals present** — named expert author, credentials, primary sources cited, clear "last updated" (ties to B6).
 - [] **F6. External corroboration exists or is being built** — the page's key facts are consistent with, and ideally referenced by, other reputable sources. *(Weakest for a brand-new site; expect AI citation to lag the Google ranking.)*
 - [] **F7. Not relying on unproven levers** — `llms.txt`, FAQ/HowTo schema, or any "guaranteed citation" trick is **not** the plan (harmless to have, but not counted on).
-

Proof artifacts to capture (for the "entry ticket")

- [] **P1. First-page ranking** — Search Console → Performance, filtered to the exact target query, showing **average position** $\leq \sim 10$. Screenshot with date.
- [] **P2. AI citation** — a screenshot of ChatGPT / Perplexity / Google AI Overview answer that links to the page. *(May take longer; not guaranteed.)*
- [] **P3. Technical health** — Rich Results Test = valid; PageSpeed/GSC CWV = all green (mobile + desktop); GSC = indexed, zero manual actions.